

RAG Eval BDD Framework

A production-grade Python framework for evaluating RAG systems using:

- pytest
- pytest-bdd (Cucumber-style Given/When/Then)
- deepeval metrics
- allure-pytest reporting
- persistent run history + trend comparison (last N runs)

This framework is isolated under `rag_eval_bdd/` and **does not modify notebooks**.

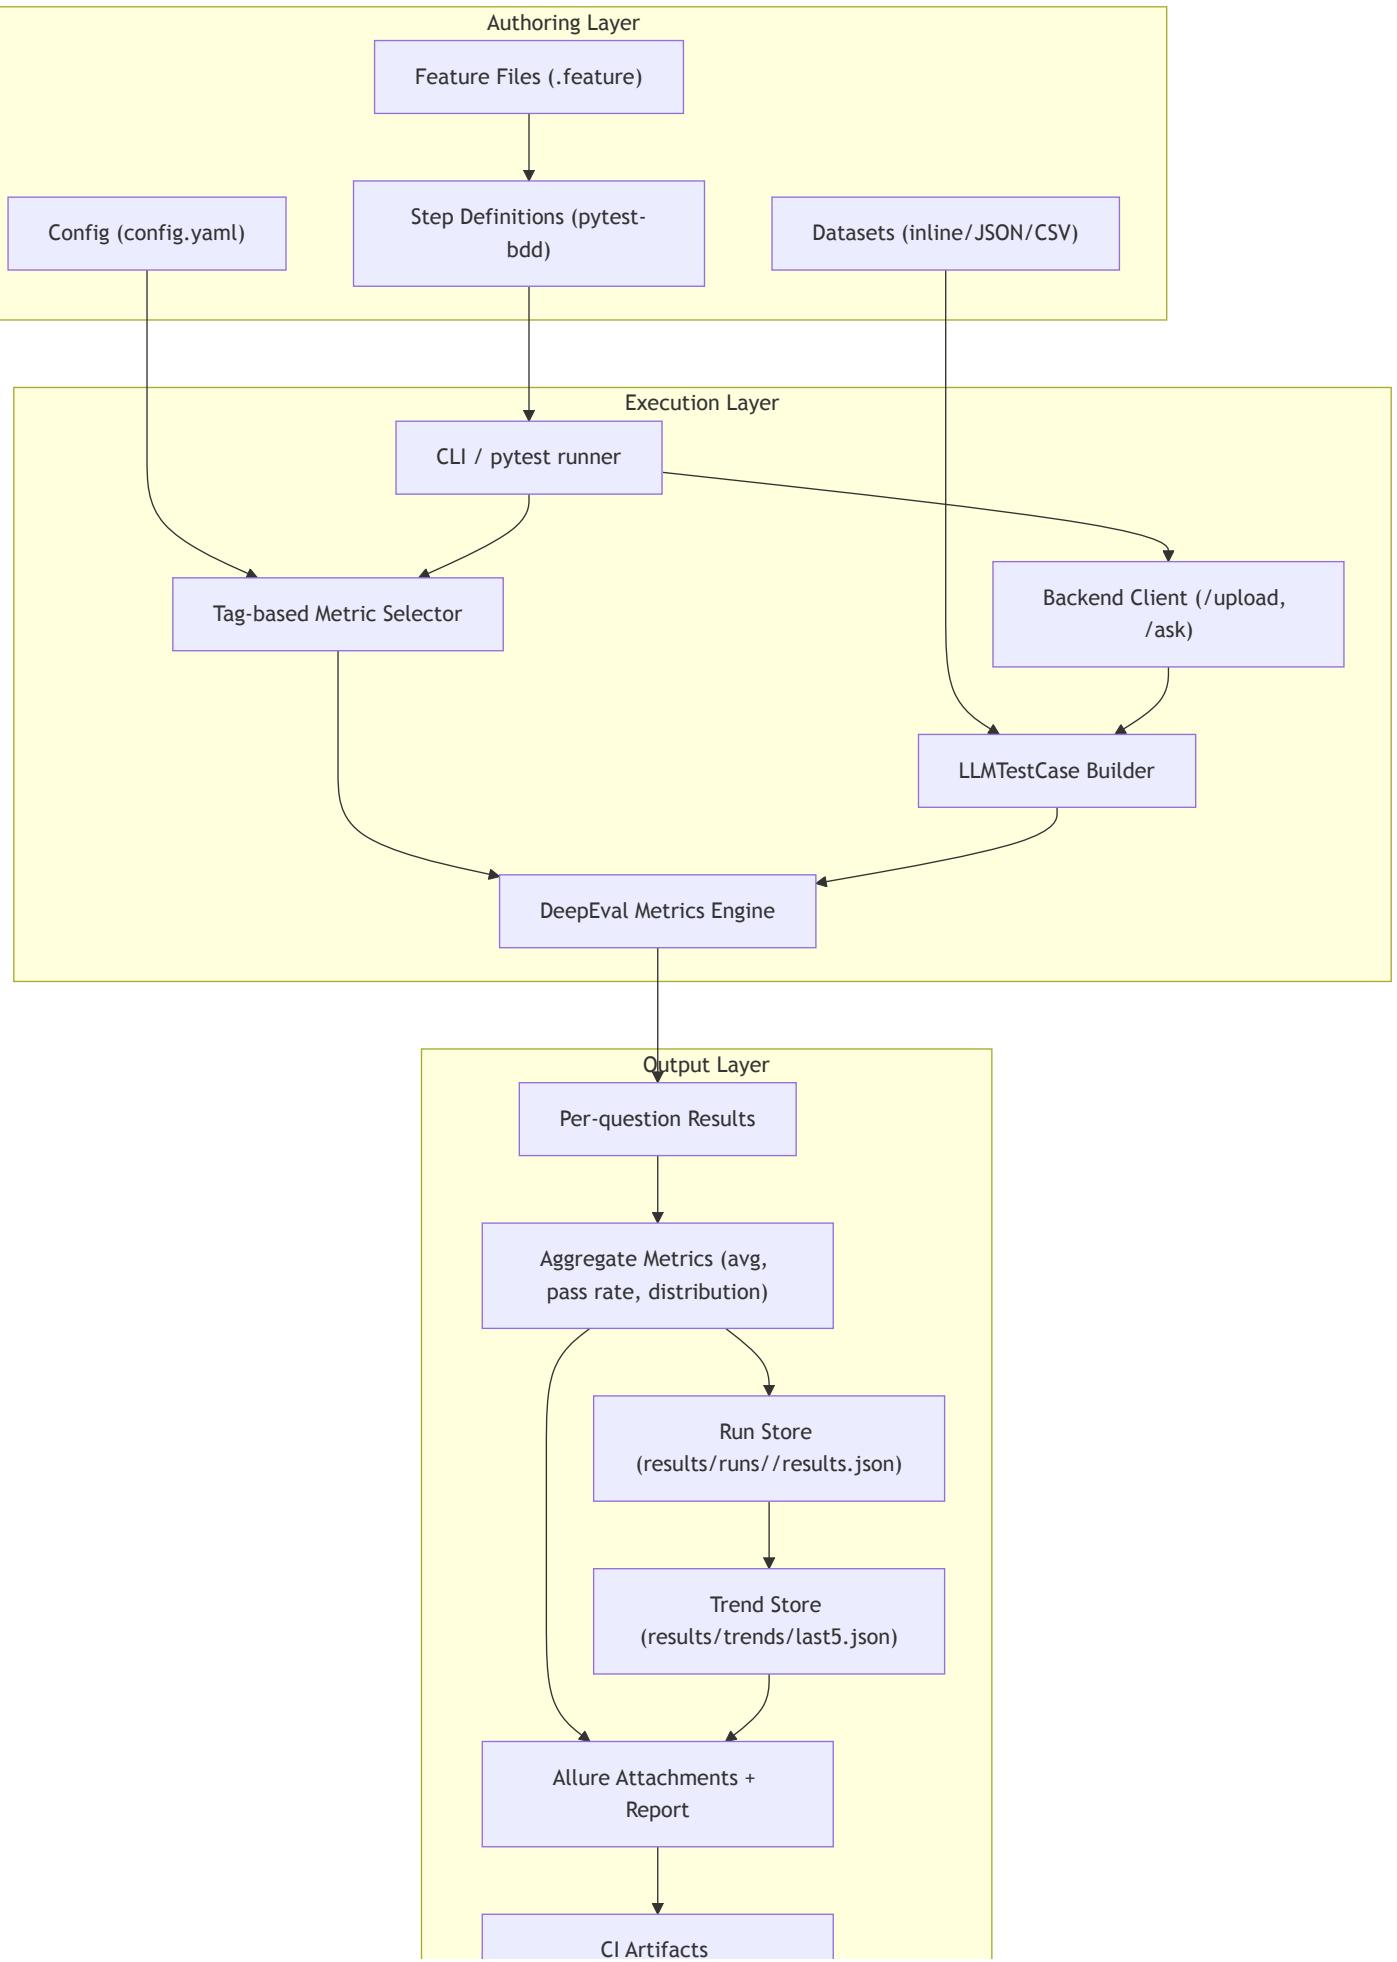
Quick Start (60 seconds)

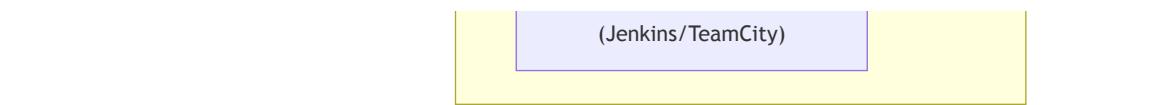
```
git clone <your-repo-url>
cd rag-session-qa-eval
source .venv/bin/activate  # reuse existing repo venv if available
cd rag_eval_bdd
python -m pip install -r requirements.txt
export BASE_URL="http://localhost:8000"
export OPENAI_API_KEY="your_openai_key"
make smoke
make live
make live-notebook-parity
```

Use `make ci` for a stable pipeline flow (`smoke` gate + non-blocking `live`).

BDD Framework Architecture (Quick Access)

If your IDE does not render Mermaid, use the ASCII architecture right below the Mermaid diagram.





AUTHORIZING LAYER

Feature Files + Step Definitions + Config + Datasets

|

v

EXECUTION LAYER

pytest/CLI -> tag-based metric selection -> backend (/upload,/ask)
-> LLMTTestCase -> DeepEval metric execution

|

v

OUTPUT LAYER

per-question results -> aggregates -> run store -> trend store
-> Allure report -> CI artifacts

1) Repository Setup

1.1 Clone the repository

```
git clone <your-repo-url>
cd rag-session-qa-eval
```

1.2 Required tools

- Python 3.11+
- pip
- venv
- Java (required by Allure CLI)
- Allure CLI (optional but recommended for local report generation)
- OpenAI API key for live DeepEval metric runs (OPENAI_API_KEY)

Check versions:

```
python3 --version
pip --version
java -version
```

1.3 Install framework dependencies

Preferred (reuse existing repo-level virtual environment):

```
cd rag_eval_bdd
source ../.venv/bin/activate
python -m pip install --upgrade pip
python -m pip install -r requirements.txt
```

If `../.venv` does not exist, create a local environment:

```
cd rag_eval_bdd
python3 -m venv .venv
source .venv/bin/activate
python -m pip install --upgrade pip
python -m pip install -r requirements.txt
```

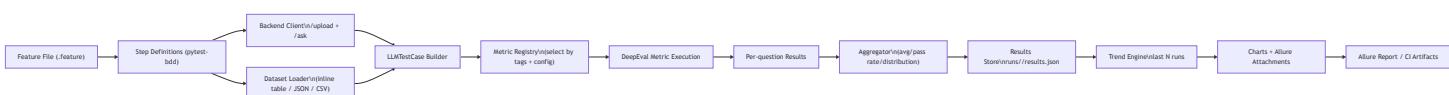
1.4 Configure environment

Set the backend URL and optional runtime settings:

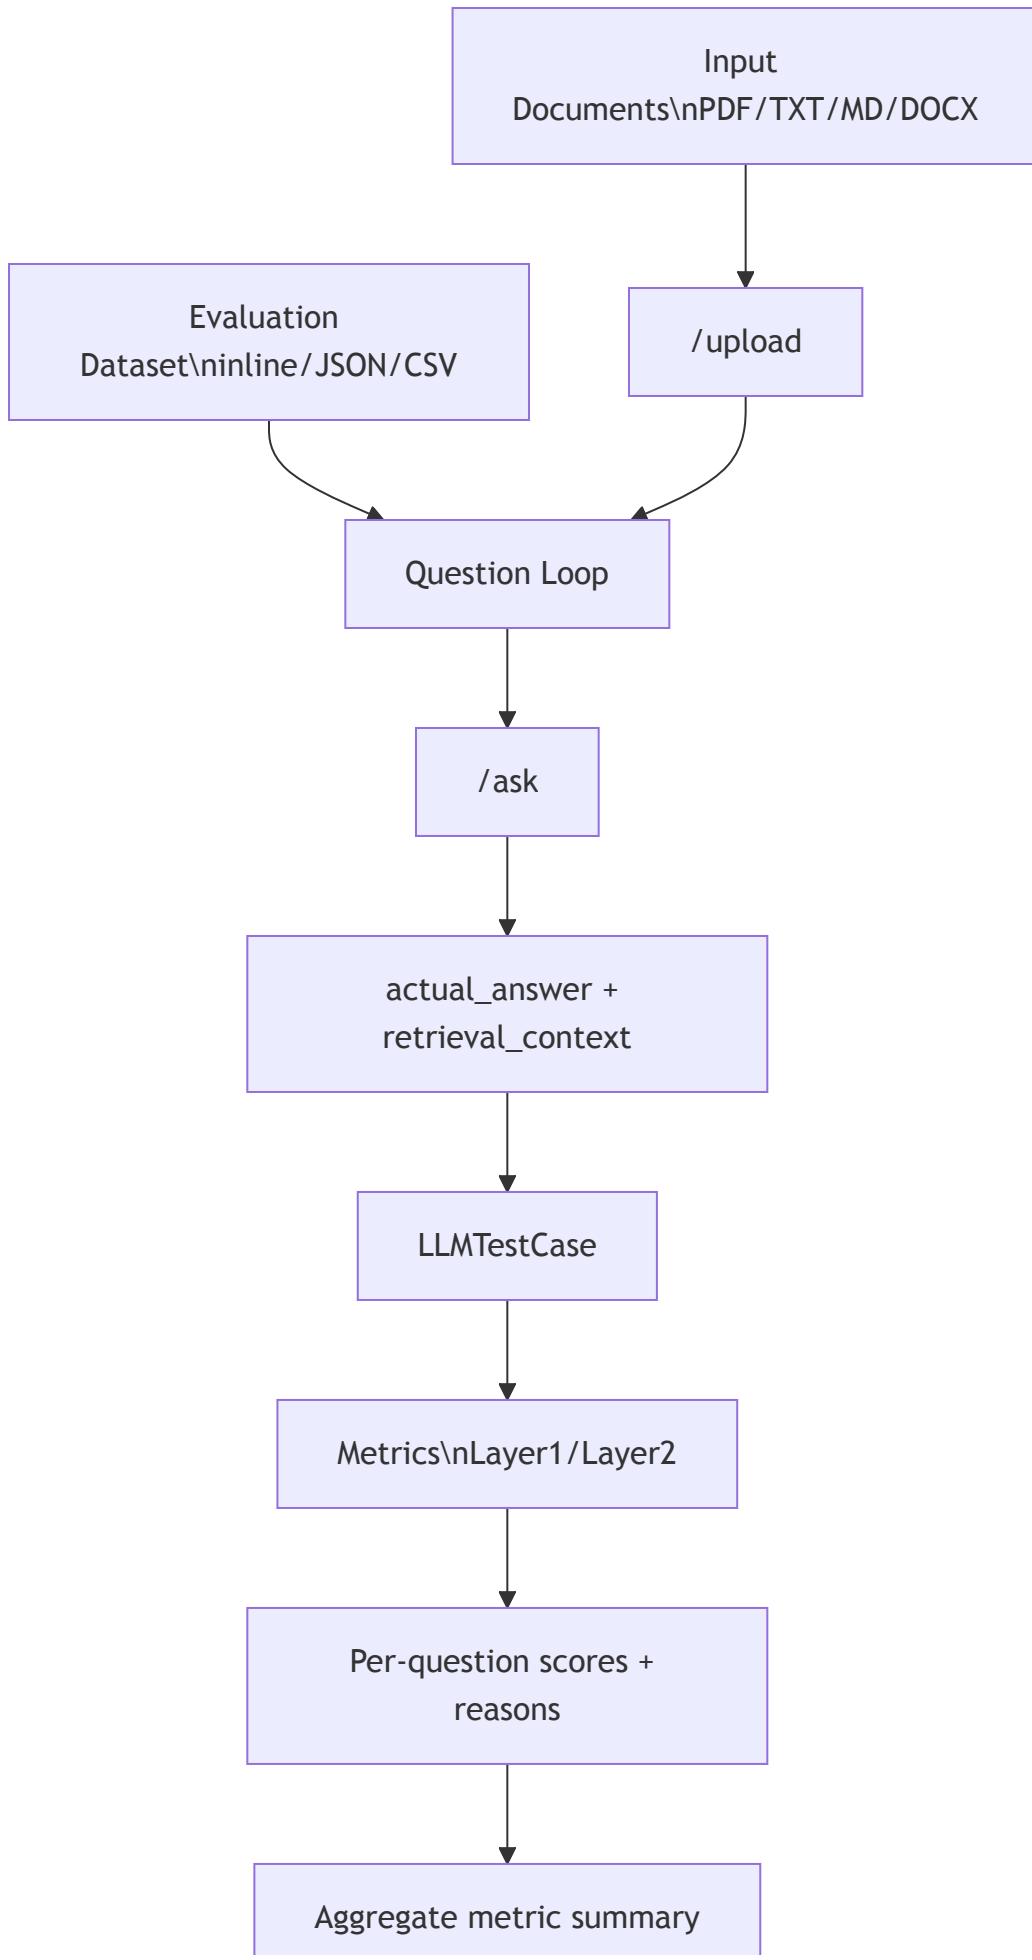
```
export BASE_URL="http://localhost:8000"
export OPENAI_API_KEY="your_openai_key" # required for live metric evaluation
export API_KEY="" # optional backend auth header
export MODEL="gpt-4.1-mini" # optional
export EMBED_MODEL="" # optional
export RAG_EVAL_CONFIG="rag_eval_bdd/config/config.yaml" # optional override
```

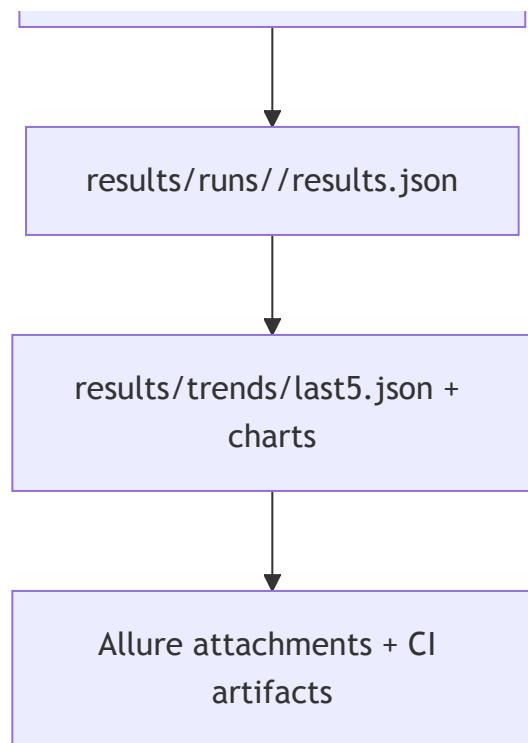
You can set these in repo-level `.env` or `rag_eval_bdd/.env`; the framework auto-loads both.

2) Workflow Diagram (End-to-End)



3) Data Flow Diagram (Input → Processing → Output)





4) Project Structure Overview

```
rag_eval_bdd/
├── config/
│   └── config.yaml
├── features/
│   ├── layer1_context_metrics.feature
│   └── layer2_answer_metrics.feature
├── steps/
│   └── test_eval_steps.py
└── src/rag_eval_bdd/
    ├── __main__.py
    ├── backend_client.py
    ├── cli.py
    ├── config_loader.py
    ├── dataset_loader.py
    ├── evaluator.py
    ├── metric_registry.py
    ├── models.py
    ├── reporting.py
    ├── results_store.py
    └── synthesize.py
├── data/
│   ├── datasets/
│   ├── generated/
│   └── source/
├── results/
│   ├── index.json
│   ├── runs/
│   └── trends/
└── tests/
    ├── test_config_loader.py
    ├── test_dataset_loader.py
    ├── test_evaluator_modes.py
    ├── test_reporting.py
    └── test_results_store.py
├── jenkins/
│   └── Jenkinsfile
├── teamcity/
│   ├── README.md
│   └── settings.kts
└── Makefile
```

```
|── pytest.ini  
└── requirements.txt
```

What goes where

- features/ : business-readable BDD scenarios.
- steps/ : step implementation code (Given/When/Then bindings).
- data/datasets/ : reusable static test datasets.
- data/generated/ : synthesized datasets (generated by CLI).
- results/runs/ : immutable run snapshots.
- results/trends/ : computed trend artifacts/charts.
- tests/ : unit tests for framework internals (non-network).

5) Configuration

Primary config file:

- rag_eval_bdd/config/config.yaml

Important sections:

- backend : base URL, endpoints, retries, timeout.
- thresholds : pass/fail criteria by metric.
- reporting : run retention and trend chart toggles.
- synthesize : default generation settings.
- evaluation : token/cost controls for metric execution.

Cost-optimized defaults

The framework is now cost-optimized by default:

- default eval model: gpt-4.1-mini
- metric reason generation disabled unless enabled
- retrieval context trimmed before metric evaluation
- document upload/session reuse across scenarios for same document
- /ask response cache reuse for repeated (session_id, question) pairs
- DeepEval retry attempts reduced to 1 (no extra retry call)

Useful overrides:

- MODEL=gpt-4.1 (or any supported model)
- RAG_EVAL_INCLUDE_REASON=1
- RAG_EVAL_MAX_CONTEXT_CHUNKS=3
- RAG_EVAL_MAX_CONTEXT_CHARS_PER_CHUNK=1200
- RAG_EVAL_DEEPEVAL_RETRY_MAX_ATTEMPTS=2
- RAG_EVAL_CACHE_UPLOADED_DOCUMENTS=0
- RAG_EVAL_CACHE_ASK_RESPONSES=0
- RAG_EVAL_FRESH_SESSION_PER_QUESTION=1
- RAG_EVAL_DISABLE_CONTEXT_TRIMMING=1
- RAG_EVAL_METRIC_QUESTION_MAPPING_MODE=all|positional|row

Notebook parity mode

Use this mode when you want live BDD behavior to be close to the notebooks:

- thresholds auto-set to 0.50
- fresh upload/session per question
- context trimming disabled
- reason text enabled
- cost optimization disabled
- positional metric-question mapping (row1->metric1, row2->metric2, row3->metric3 when counts match)

Enable with:

- RAG_EVAL_NOTEBOOK_PARITY_MODE=1
- or make live-notebook-parity

Threshold policy

Thresholds are **only** in config, not in feature files. This keeps business scenarios clean and allows centralized quality control.

6) Test Case Guidelines

6.1 Writing new test cases

Use either:

1. Inline dataset tables in feature files.
2. External dataset files (JSON/CSV) loaded by a step.

6.2 Naming conventions

- Feature file: <layer>_<intent>.feature
 - Example: layer2_answer_metrics.feature
- Scenario names: explicit business intent
 - Example: Evaluate layer2 answer metrics from external dataset file
- Dataset ids: stable and traceable
 - Example: L2_EXT_001

6.3 Best practices

- Keep one scenario focused on one evaluation objective.
- Reuse dataset files for stable regression coverage.
- Add category/topic fields to slice failures quickly.
- Keep expected answers concise and factual.
- Avoid overfitting thresholds to one run.

7) Feature File Creation (Gherkin)

7.1 Minimal template

```
@layer2 @answer_relevancy @faithfulness
Feature: Layer 2 answer quality
```

```
Scenario: Evaluate answer metrics
```

```
  Given backend is reachable
```

```
  And documents are uploaded from "eval/sample_docs/Match_Summary.pdf"
```

```
  And I load dataset "rag_eval_bdd/data/datasets/layer2_questions.json"
```

```
  When I evaluate all questions
```

```
  Then metric "answer_relevancy" should be >= configured threshold
```

```
  And metric "faithfulness" should be >= configured threshold
```

```
  And save results for reporting
```

7.2 Inline table example

```
And I use inline dataset:
```

id	question	expected_answer	
Q001	How many sixes did Tilak Varma hit?	Tilak Varma hit 3 sixes.	ba
Q002	Who dismissed Suryakumar Yadav?	Kwena Maphaka dismissed Surya.	wi

7.3 Readability and reuse tips

- Keep Given steps generic and reusable.
- Put large datasets in files, not inline.
- Use tags to control metric scope.
- Keep scenarios short; move heavy details to data files.

8) Step Definitions

Step bindings are in:

- `rag_eval_bdd/steps/test_eval_steps.py`

8.1 Binding pattern

- `@given(...)` prepares backend/session/data.
- `@when(...)` executes evaluation loop.
- `@then(...)` asserts aggregate thresholds and persists artifacts.

8.2 Typical organization pattern

- Input setup steps
- Execution step(s)
- Assertion steps
- Persistence/reporting step

8.3 New step template

```
from pytest_bdd import given, when, then, parsers

@given(parsers.parse('I set category "{category}" filter'))
def set_category_filter(category, scenario_state):
    scenario_state.category_filter = category

@when("I evaluate filtered questions")
def evaluate_filtered(...):
    ...

@then(parsers.parse('metric "{metric_name}" pass rate should be >= {min_rate:f}%'))
def assert_pass_rate(metric_name, min_rate, scenario_state):
    ...
```

9) Dataset Creation and Management

9.1 Supported formats

- JSON (`.json`)
- CSV (`.csv`)
- text lines (`.txt` , `.md`) for basic loading/synthesis context extraction

9.2 Recommended dataset schema

```
[  
  {  
    "id": "L2_EXT_001",  
    "question": "How many sixes did Tilak Varma hit?",  
    "expected_answer": "Tilak Varma hit 3 sixes.",  
    "category": "batting",  
    "source_reference": "Match_Summary.pdf"  
  }  
]
```

9.3 CSV schema

```
id,question,expected_answer,category,source_reference  
L2_EXT_001,How many sixes did Tilak Varma hit?,Tilak Varma hit 3 sixes.,batting,Match_S
```

9.4 Data reuse guidelines

- Keep canonical datasets under `data/datasets/`.
- Keep generated datasets under `data/generated/`.
- Use stable IDs for trend comparison and triage.
- Version control curated datasets.

10) Synthetic Dataset Generation (DeepEval Synthesizer)

Generate from folder/file/dataset:

```
cd rag_eval_bdd  
python -m rag_eval_bdd synthesize \  
  --input data/source/ \  
  --output data/generated/questions.json \  
  --num-questions 50
```

Generated output fields:

- id
- question
- expected_answer (optional)
- category (optional)
- source_reference (optional)

11) Execution and Reports

11.1 Test execution

Note: run pytest against steps/ (not features/) because pytest-bdd scenarios are materialized from Python step modules.

Recommended (via Makefile):

```
cd rag_eval_bdd
make smoke          # deterministic and stable
make live           # live backend + LLM metrics
make live-notebook-parity # notebook-like behavior for live demo parity
make ci              # smoke gate + live non-blocking
```

Deterministic smoke suite (stable CI gate, no live LLM/backend dependency):

```
cd rag_eval_bdd
python -m pytest -c pytest.ini tests -m "smoke" -q
```

Live BDD evaluation suite:

```
python -m pytest -c pytest.ini steps -m "live" --alluredir=allure-results
```

Notebook-parity live run:

```
RAG_EVAL_NOTEBOOK_PARITY_MODE=1 python -m pytest -c pytest.ini steps -m "live" --allure
```

Run all live features:

```
python -m pytest -c pytest.ini steps --alluredir=allure-results
```

Live suite as non-blocking quality signal:

```
python -m pytest -c pytest.ini steps -m "live" --alluredir=allure-results || true
```

By default, the framework now does this automatically for every run (when `--alluredir` is present):

- clears old `allure-results/` and `allure-report/`
- generates a fresh `allure-report/`
- opens the report in browser via `allure open` (local HTTP server)

Run layer-specific:

```
python -m pytest -c pytest.ini steps -m "layer1" --alluredir=allure-results  
python -m pytest -c pytest.ini steps -m "layer2" --alluredir=allure-results
```

Run metric subset:

```
python -m pytest -c pytest.ini steps -m "faithfulness or completeness" --alluredir=alu
```

Run via CLI wrapper:

```
python -m rag_eval_bdd run --suite smoke  
python -m rag_eval_bdd run --suite live --tags "layer1"  
python -m rag_eval_bdd run --suite live --feature features/layer2_answer_metrics.feature
```

`--feature` currently derives layer filter (`layer1` / `layer2`) from the feature file name.

11.2 Allure report generation

Auto-generation happens at the end of each pytest run. Manual command is still available:

```
allure generate allure-results --clean -o allure-report
```

11.3 Output meaning

- `results/runs/<run_id>/results.json`

- full run payload (question-level + aggregate)
- results/index.json
 - pointers to recent runs
- results/trends/last5.json
 - trend summary across recent runs
- Allure attachments
 - per-question metrics
 - raw request/response
 - trend charts and trend HTML summary

12) CI/CD Integration

12.1 Jenkins

Use provided pipeline:

- rag_eval_bdd/jenkins/Jenkinsfile

It performs:

1. Checkout
2. Create virtualenv
3. Install dependencies
4. Run deterministic smoke gate
5. Run live eval as non-blocking stage
6. Publish Allure results
7. Archive results/ artifacts

12.2 GitHub Actions (example)

Create .github/workflows/rag-eval-bdd.yml :

```
name: rag-eval-bdd

on:
  workflow_dispatch:
  push:
    paths:
      - "rag_eval_bdd/**"

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-python@v5
        with:
          python-version: "3.11"
      - name: Install
        run: |
          cd rag_eval_bdd
          python -m venv .venv-ci
          . .venv-ci/bin/activate
          pip install --upgrade pip
          pip install -r requirements.txt
      - name: Run tests
        env:
          BASE_URL: ${{ secrets.BASE_URL }}
          OPENAI_API_KEY: ${{ secrets.OPENAI_API_KEY }}
          API_KEY: ${{ secrets.API_KEY }}
        run: |
          cd rag_eval_bdd
          . .venv-ci/bin/activate
          pytest -c pytest.ini tests -m "smoke" -q
          pytest -c pytest.ini steps -m "live and (layer1 or layer2)" --alluredir=allure
      - name: Upload artifacts
        uses: actions/upload-artifact@v4
        with:
          name: rag-eval-artifacts
          path: |
            rag_eval_bdd/allure-results
            rag_eval_bdd/results
```

12.3 TeamCity

See:

- `rag_eval_bdd/teamcity/README.md`
- `rag_eval_bdd/teamcity/settings.kts`

13) Troubleshooting

Error: ModuleNotFoundError: No module named 'pytest_bdd'

Fix:

```
cd rag_eval_bdd
source ../../venv/bin/activate
python -m pip install -r requirements.txt
```

Error: no tests ran when using pytest ... features

Cause:

- `pytest-bdd` scenarios are materialized from Python step modules, so collecting only `features/` does not execute scenarios.

Fix:

```
python -m pytest -c pytest.ini steps --alluredir=allure-results
```

Error:

`pytest: error: unrecognized arguments: --alluredir=...`

Cause:

- `allure-pytest` plugin is not installed in the active virtual environment.

Fix:

```
cd rag_eval_bdd
source ../.venv/bin/activate
python -m pip install -r requirements.txt
python -m pytest -c pytest.ini steps --alluredir=allure-results
```

Error: Backend not reachable

Symptoms:

- `RuntimeError: Backend is not reachable`

Checks:

- Backend process is running
- `BASE_URL` is correct
- network/firewall allows access

Error: /upload or /ask failures

Checks:

- endpoint paths in `config/config.yaml`
- request payload compatibility with backend
- API auth (`API_KEY`) if required

Error: No metrics selected

Cause:

- tag selection does not map to layer/metric tags

Fix:

- include `@layer1` or `@layer2`
- or include explicit metric tags (`@faithfulness` , etc.)

Error: Completeness/Context relevance class unavailable

Behavior:

- framework falls back to `GEval` for unsupported metric classes

Flaky score movements across runs

Mitigation:

- evaluate using distributions and pass rate, not only single score
- compare against trend (`last5.json`)
- tune thresholds conservatively in config

Allure report empty or not rendering

Checks:

- tests executed with `--alluredir=allure-results`
- allure CLI installed and Java available
- do not open `allure-report/index.html` directly with `file://` in browser; use `allure open allure-report`
- if browser shows endless "Loading...", rerun tests and let auto-open handle it (served over HTTP)

Auto-report controls:

- `RAG_EVAL_AUTO_ALLURE_CLEAN=0` to keep previous `allure-results`
- `RAG_EVAL_AUTO_ALLURE_GENERATE=0` to disable auto generation
- `RAG_EVAL_AUTO_ALLURE_OPEN=0` to disable browser popup

14) Quick Start (Minimal)

```
cd rag_eval_bdd
source ../.venv/bin/activate
python -m pip install -r requirements.txt
export BASE_URL="http://localhost:8000"
export OPENAI_API_KEY="your_openai_key"
make live
```

15) Contributor Checklist

Before pushing changes:

1. Add/update tests in `rag_eval_bdd/tests/` for non-network logic.
2. Keep feature files business-readable.
3. Do not hardcode thresholds in feature files.
4. Ensure new steps are reusable and generic.
5. Verify outputs are persisted and trend-compatible.
6. Keep notebooks untouched.